


Transfer learning for text classification

DEEP LEARNING FOR TEXT WITH PYTORCH



Shubham Jain
Instructor

What is transfer learning?



- Use pre-existing knowledge from one task to a related task
- Saves time
- Share expertise
- Reduces need for large data
- An English teacher starts teaching History

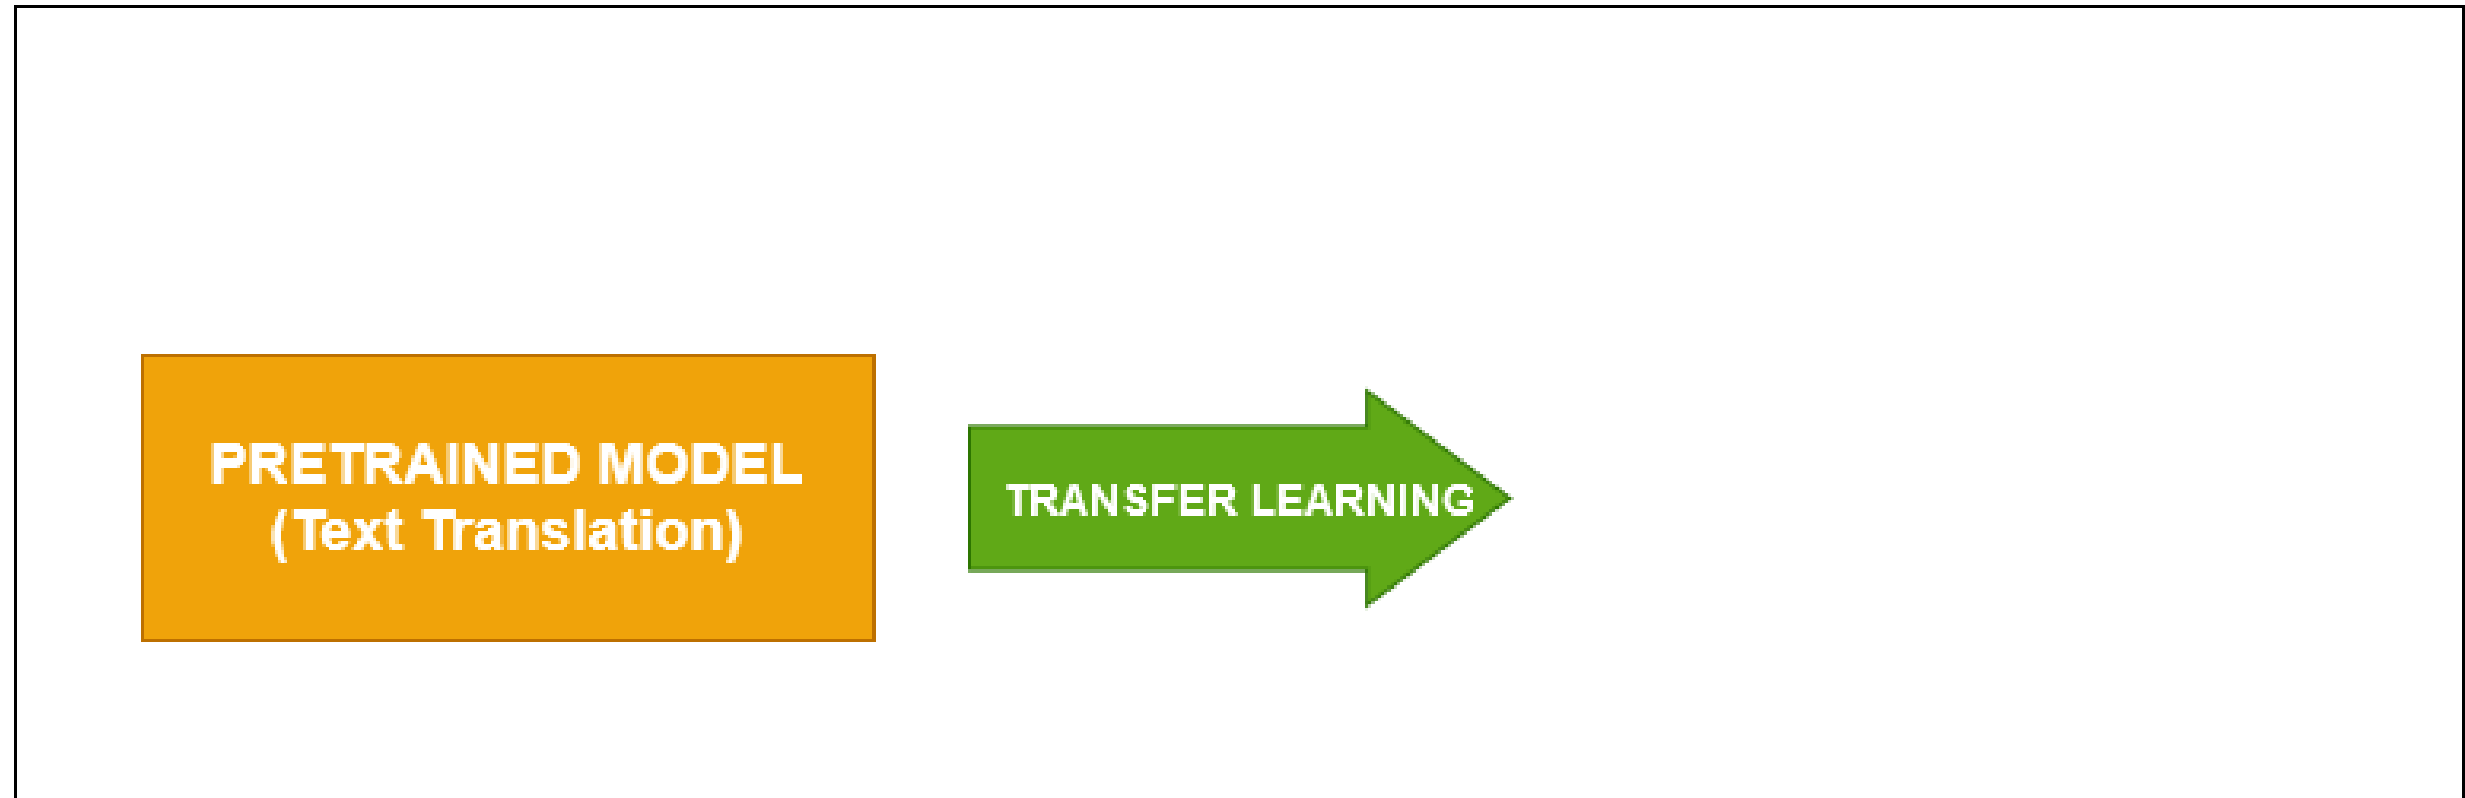
Mechanics of transfer learning



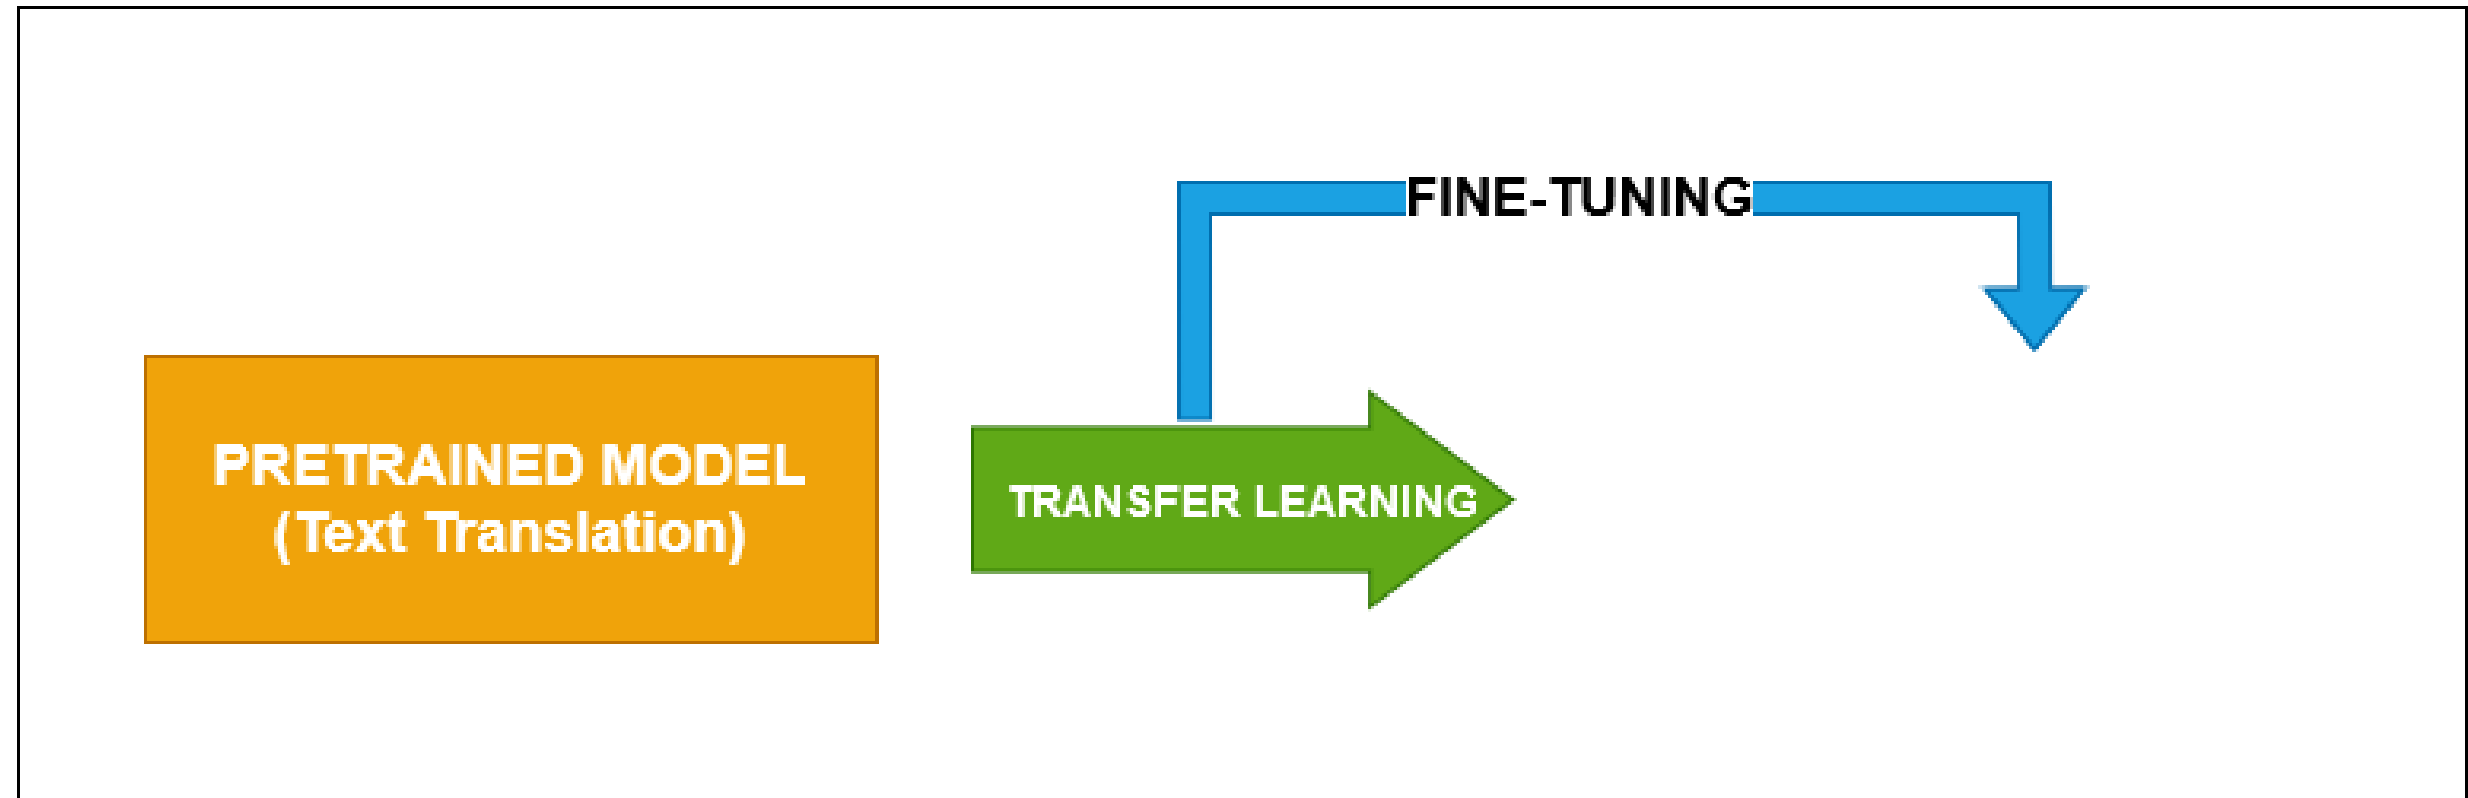
A diagram illustrating the mechanics of transfer learning. It features a large, empty rectangular box with a thin black border. Inside this box, on the left side, is a smaller orange rectangle. The orange rectangle contains the text "PRETRAINED MODEL (Text Translation)" in white, bold, uppercase letters.

PRETRAINED MODEL
(Text Translation)

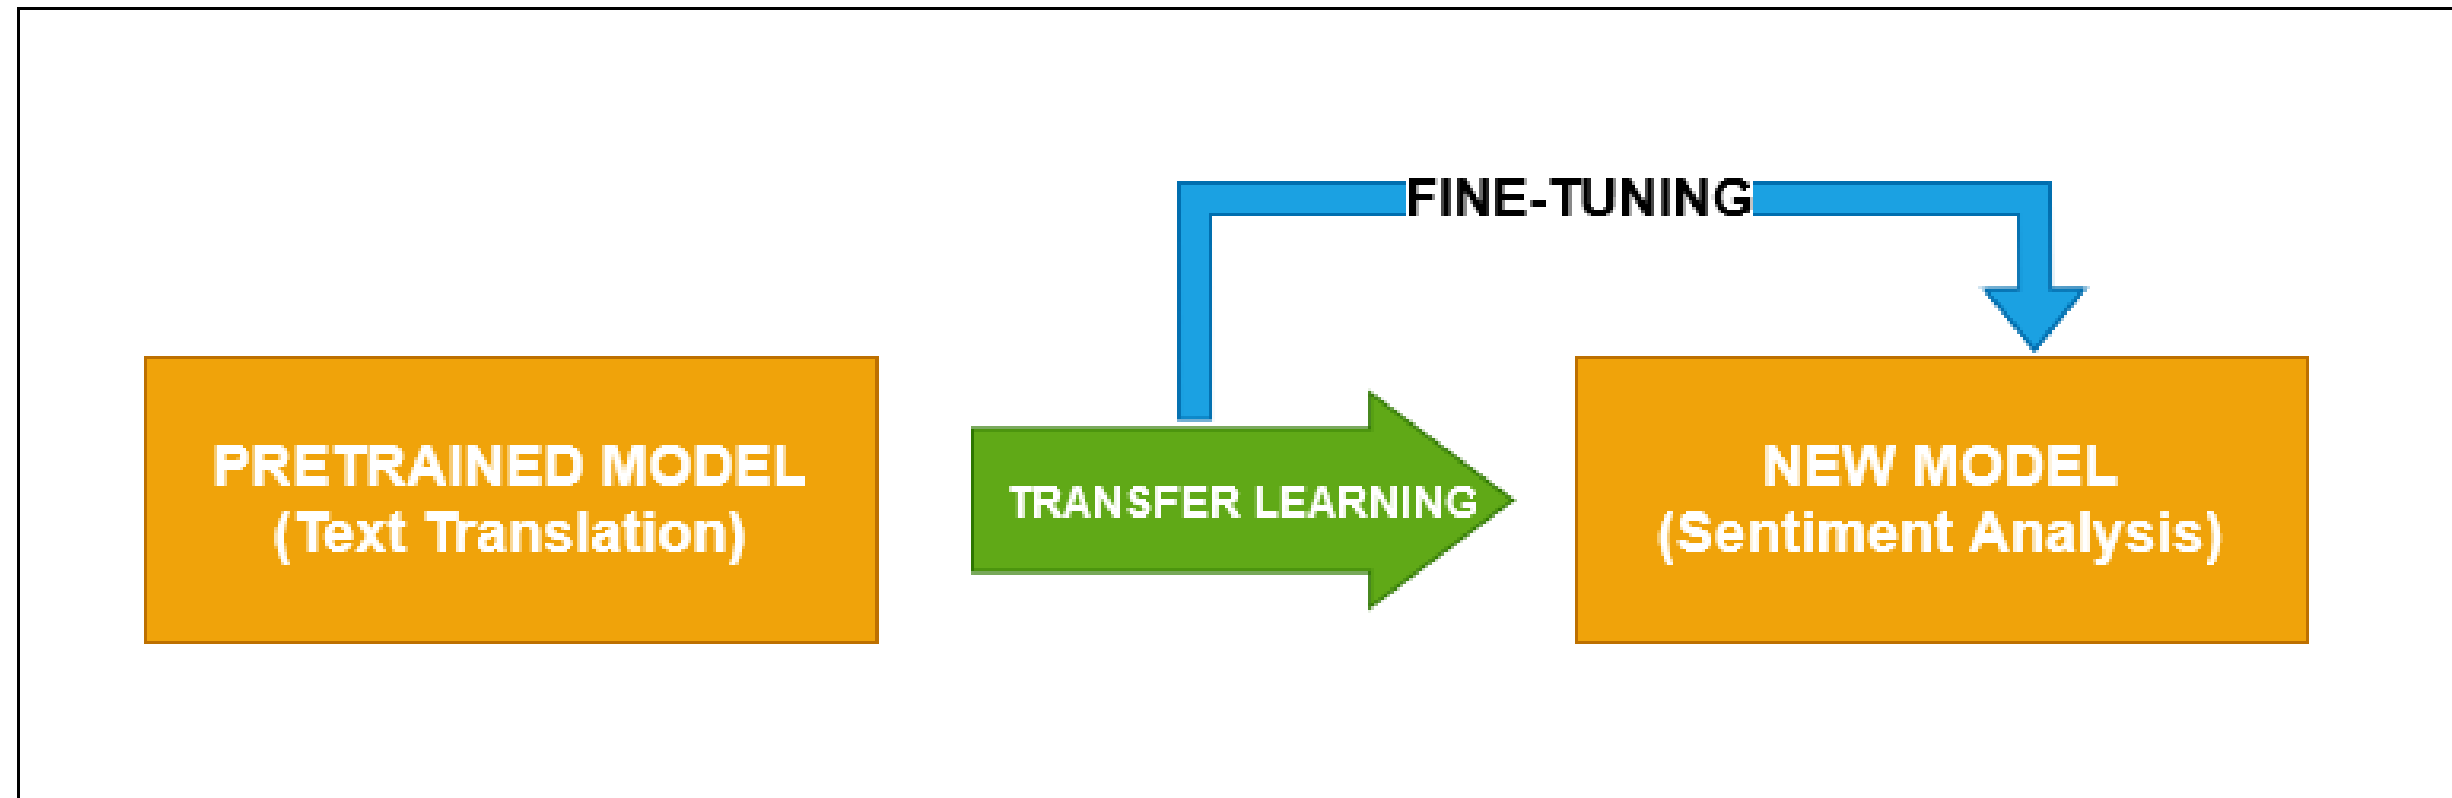
Mechanics of transfer learning



Mechanics of transfer learning




Mechanics of transfer learning



Pre-trained model : BERT

- Bidirectional Encoder Representations from Transformers



- Trained for language modeling
- Multiple layers of transformers
- Pre-trained on large texts

Hands-on: implementing BERT

```
texts = ["I love this!",
         "This is terrible.",
         "Amazing experience!",
         "Not my cup of tea."]
labels = [1, 0, 1, 0]
import torch
from transformers import BertTokenizer, BertForSequenceClassification
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
                                                    num_labels=2)
inputs = tokenizer(texts, padding=True, truncation=True,
                  return_tensors="pt", max_length=32)
inputs["labels"] = torch.tensor(labels)
```


Fine-tuning BERT

```
optimizer = torch.optim.AdamW(model.parameters(), lr=0.00001)
model.train()
for epoch in range(1):
    outputs = model(**inputs)
    loss = outputs.loss
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()
    print(f"Epoch: {epoch+1}, Loss: {loss.item()}")
```

```
Epoch: 1, Loss: 0.7061821222305298
```

Evaluating on new text

```
text = "I had an awesome day!"
input_eval = tokenizer(text, return_tensors="pt", truncation=True,
                       padding=True, max_length=128)
outputs_eval = model(**input_eval)
predictions = torch.nn.functional.softmax(outputs_eval.logits, dim=-1)
predicted_label = 'positive' if torch.argmax(predictions) > 0 else 'negative'
print(f"Text: {text}\nSentiment: {predicted_label}")
```

```
Text: I had an awesome day!
Sentiment: positive
```

Let's practice!

DEEP LEARNING FOR TEXT WITH PYTORCH

Transformers for text processing

DEEP LEARNING FOR TEXT WITH PYTORCH



Shubham Jain
Instructor

Why use transformers for text processing?



Transformers

- Speed
- Understand the relationship between words, regardless of distances
- Human-like response

Components of a transformer

- **Encoder:** Processes input data
- **Decoder:** Reconstructs the output
- **Feed-forward Neural Networks:** Refine understanding
- **Positional Encoding:** Ensure order matters
- **Multi-Head Attention:** Captures multiple inputs or sentiments

Preparing our data: train-test split

```
sentences = ["I love this product", "This is terrible",  
             "Could be better", "This is the best"]  
  
labels = [1, 0, 0, 1]  
  
train_sentences = sentences[:3]  
train_labels = labels[:3]  
test_sentences = sentences[3:]  
test_labels = labels[3:]
```

Building the transformer model

```
class TransformerEncoder(nn.Module):
    def __init__(self, embed_size, heads, num_layers, dropout):
        super(TransformerEncoder, self).__init__()
        self.encoder = nn.TransformerEncoder(
            nn.TransformerEncoderLayer(d_model=embed_size, nhead=heads),
            num_layers=num_layers)
        self.fc = nn.Linear(embed_size, 2)

    def forward(self, x):
        x = self.encoder(x)
        x = x.mean(dim=1)
        return self.fc(x)

model = TransformerEncoder(embed_size=512, heads=8, num_layers=3, dropout=0.5)
optimizer = optim.Adam(model.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()
```


Training the transformers

```
for epoch in range(5):
    for sentence, label in zip(train_sentences, train_labels):
        tokens = sentence.split()
        data = torch.stack([token_embeddings[token] for token in tokens], dim=1)
        output = model(data)
        loss = criterion(output, torch.tensor([label]))
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
    print(f"Epoch {epoch}, Loss: {loss.item()}")
```

```
Epoch 0, Loss: 13.788233757019043
Epoch 1, Loss: 3.9480819702148438
Epoch 2, Loss: 2.4790847301483154
Epoch 3, Loss: 1.3020926713943481
Epoch 4, Loss: 0.4660853147506714
```

Predicting the transformers

```
def predict(sentence):  
    model.eval()  
    with torch.no_grad():  
        tokens = sentence.split()  
        data = torch.stack([token_embeddings.get(token, torch.rand((1, 512)))  
                             for token in tokens], dim=1)  
        output = model(data)  
        predicted = torch.argmax(output, dim=1)  
        return "Positive" if predicted.item() == 1 else "Negative"
```

Predicting on new text

```
sample_sentence = "This product can be better"  
print(f"'{sample_sentence}' is {predict(sample_sentence)}")
```

```
'This product can be better' is Negative
```

Let's practice!

DEEP LEARNING FOR TEXT WITH PYTORCH

Attention mechanisms for text generation

DEEP LEARNING FOR TEXT WITH PYTORCH



Shubham Jain
Instructor

The ambiguity in text processing

- "The monkey ate that banana because it was too hungry"
- What does the word "it" refer to?

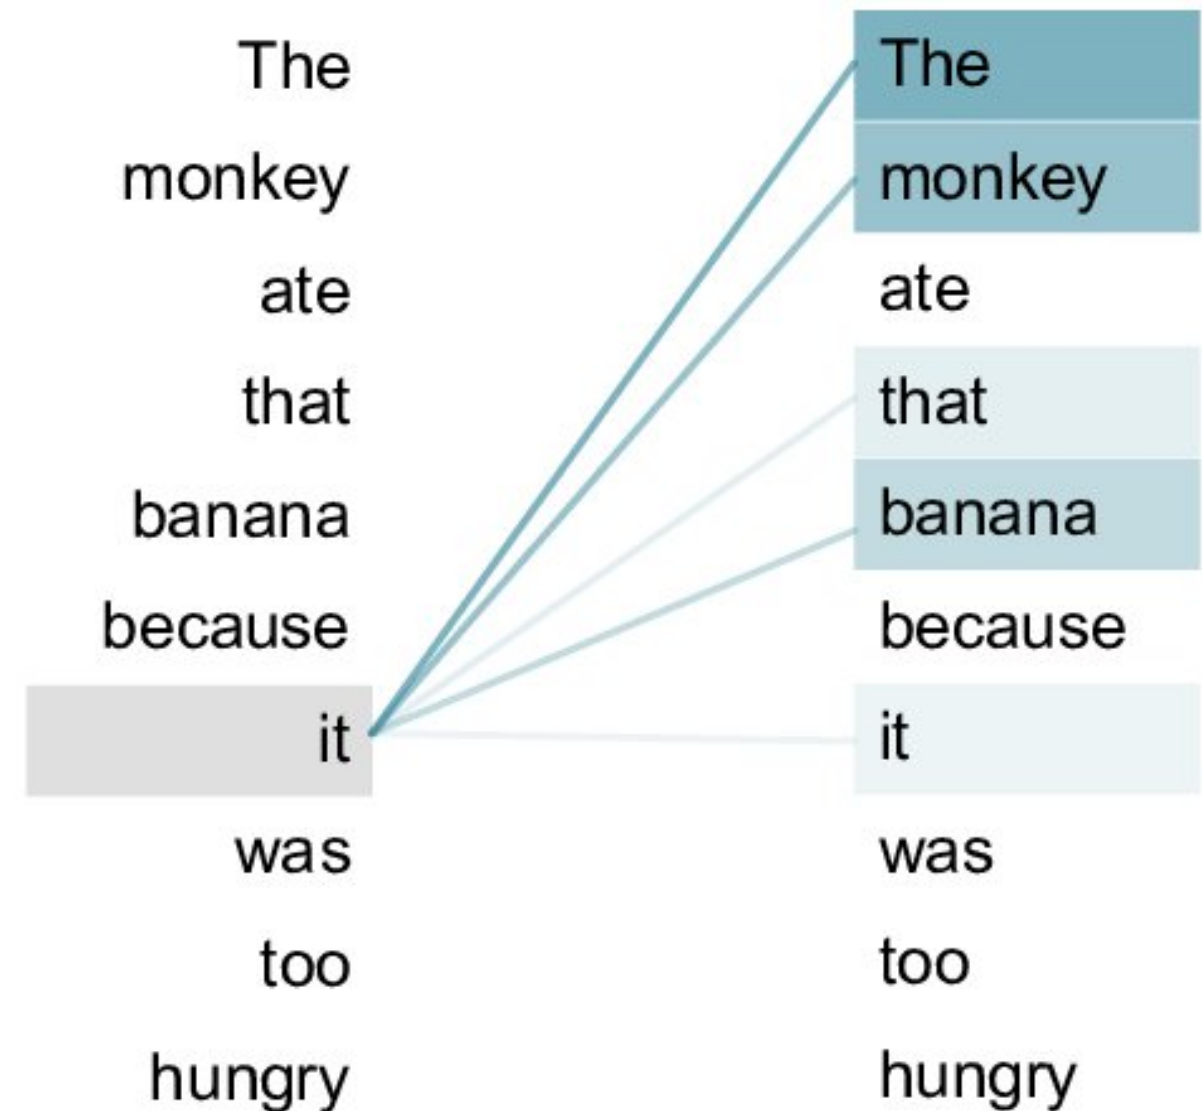


IT?



Attention mechanisms

- Assigns importance to words
- Ensures that machine's interpretation aligns with human understanding



¹ Xie, Huiqiang & Qin, Zhijin & Li, Geoffrey & Juang, Biing-Hwang. (2020). Deep Learning Enabled Semantic Communication Systems

Self and multi-head attention

- Self-Attention: assigns significance to words within a sentence
 - "The cat, which was on the roof, was scared"
 - Linking "was scared" to "The cat"
- Multi-Head Attention: like having multiple spotlights, capturing different facets
 - Understanding "was scared" can relate to
 - "The cat", "the roof", or "was on"

Attention mechanism - setting vocabulary and data

```
data = ["the cat sat on the mat", ...]
vocab = set(' '.join(data).split())
word_to_ix = {word: i for i, word in enumerate(vocab)}
ix_to_word = {i: word for word, i in word_to_ix.items()}
pairs = [sentence.split() for sentence in data]
input_data = [[word_to_ix[word] for word in sentence[:-1]] for sentence in pairs]
target_data = [word_to_ix[sentence[-1]] for sentence in pairs]
inputs = [torch.tensor(seq, dtype=torch.long) for seq in input_data]
targets = torch.tensor(target_data, dtype=torch.long)
```

Model definition

```
embedding_dim = 10
hidden_dim = 16

class RNNWithAttentionModel(nn.Module):
    def __init__(self):
        super(RNNWithAttentionModel, self).__init__()
        self.embeddings = nn.Embedding(vocab_size, embedding_dim)
        self.rnn = nn.RNN(embedding_dim, hidden_dim, batch_first=True)
        self.attention = nn.Linear(hidden_dim, 1)
        self.fc = nn.Linear(hidden_dim, vocab_size)
```

Forward propagation with attention

```
def forward(self, x):
    x = self.embeddings(x)
    out, _ = self.rnn(x)
    attn_weights = torch.nn.functional.softmax(self.attention(out).squeeze(2),
                                                dim=1)
    context = torch.sum(attn_weights.unsqueeze(2) * out, dim=1)
    out = self.fc(context)
    return out

def pad_sequences(batch):
    max_len = max([len(seq) for seq in batch])
    return torch.stack([torch.cat([seq, torch.zeros(max_len-len(seq)).long()])
                        for seq in batch])
```

Training preparation

```
criterion = nn.CrossEntropyLoss()
attention_model = RNNWithAttentionModel()
optimizer = torch.optim.Adam(attention_model.parameters(), lr=0.01)
for epoch in range(300):
    attention_model.train()
    optimizer.zero_grad()
    padded_inputs = pad_sequences(inputs)
    outputs = attention_model(padded_inputs)
    loss = criterion(outputs, targets)
    loss.backward()
    optimizer.step()
```

Model evaluation

```
for input_seq, target in zip(input_data, target_data):
    input_test = torch.tensor(input_seq, dtype=torch.long).unsqueeze(0)
    attention_model.eval()
    attention_output = attention_model(input_test)
    attention_prediction = ix_to_word[torch.argmax(attention_output).item()]
    print(f"\nInput: {' '.join([ix_to_word[ix] for ix in input_seq])}")
    print(f"Target: {ix_to_word[target]}")
    print(f"RNN with Attention prediction: {attention_prediction}")
```

Input: the cat sat on the

Target: mat

RNN with Attention prediction: mat

Let's practice!

DEEP LEARNING FOR TEXT WITH PYTORCH

Adversarial attacks on text classification models

DEEP LEARNING FOR TEXT WITH PYTORCH



Shubham Jain
Instructor

What are adversarial attacks?

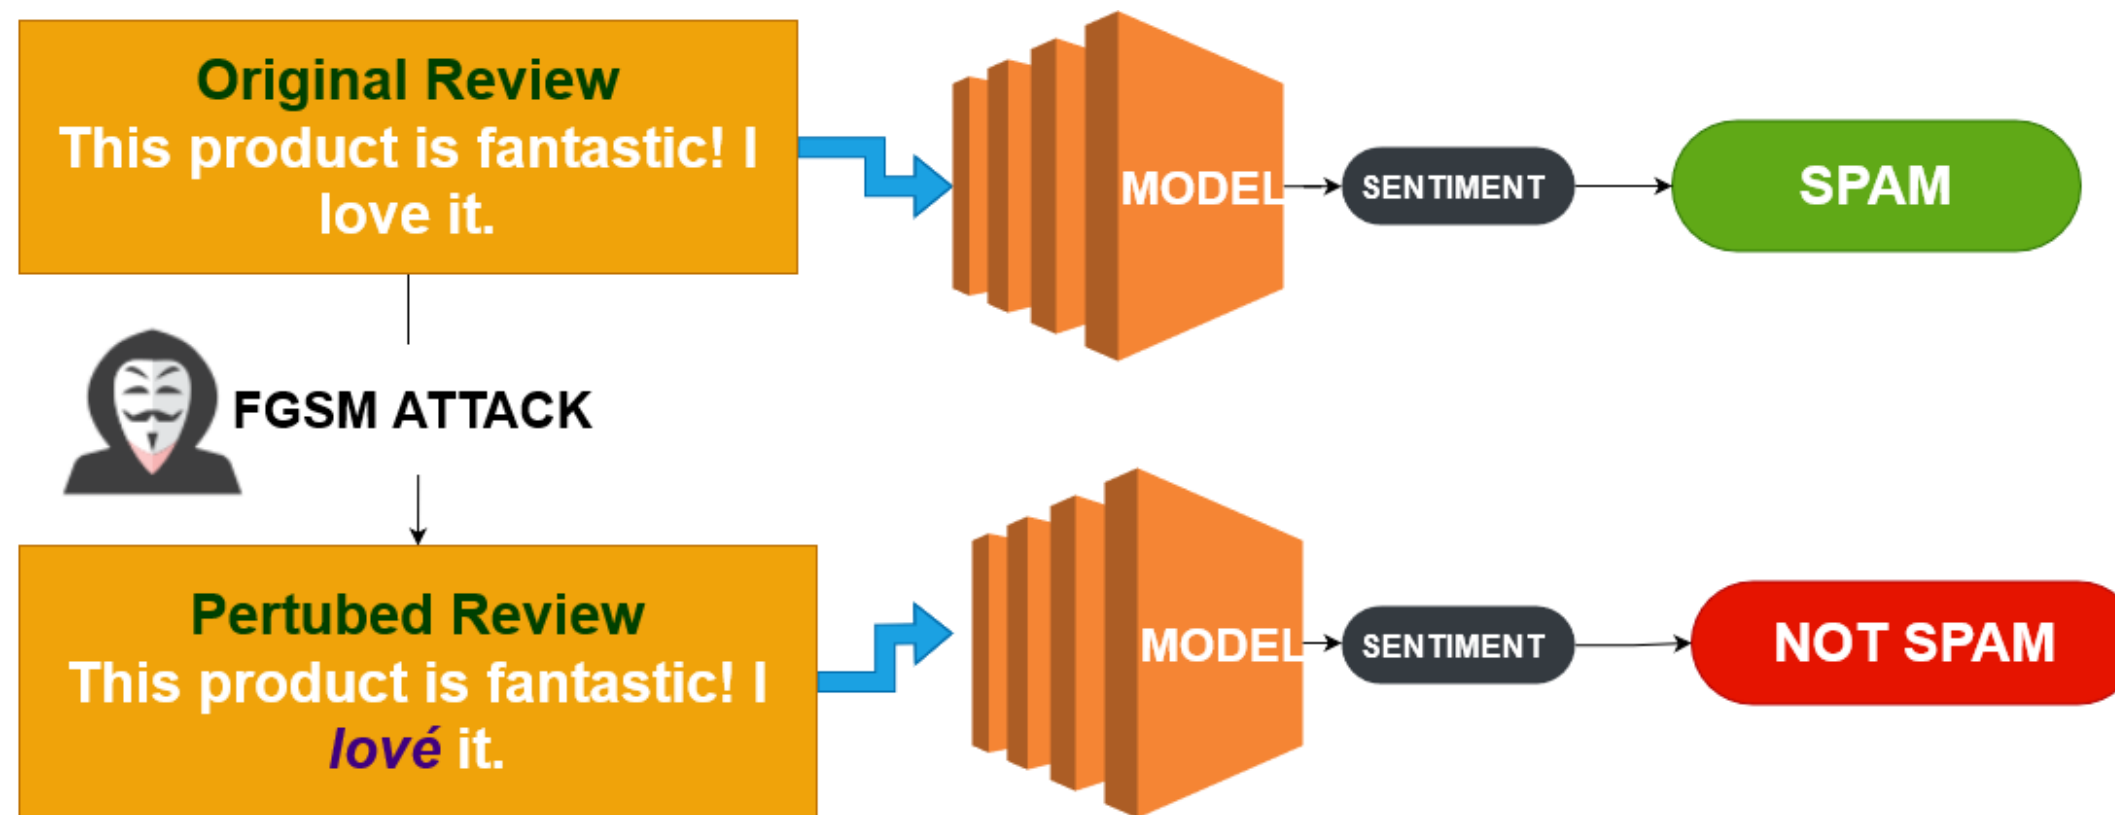
- Tweaks to input data
- Not random but calculated malicious changes
- Can drastically affect AI's decision-making

Importance of robustness

- AI systems deciding if user comments are toxic or benign
- AI unintentionally amplifying negative stereotypes from biased data
- AI giving misleading information

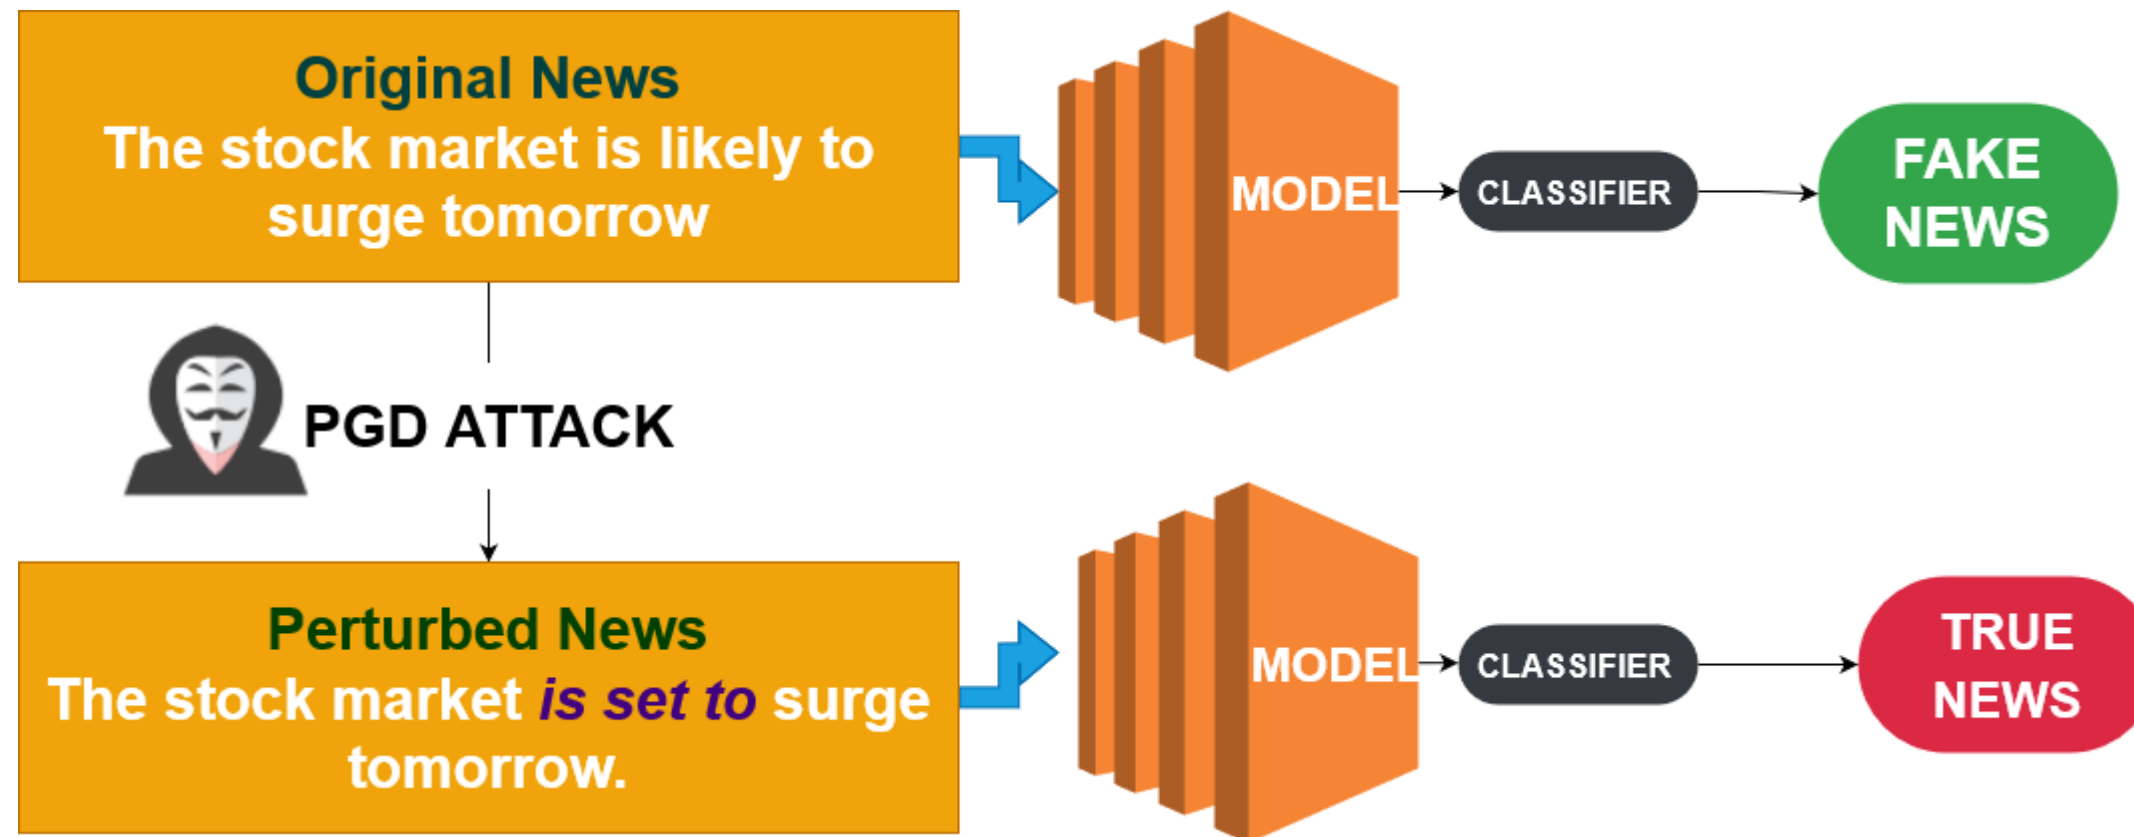
Fast Gradient Sign Method (FGSM)

- Exploits the model's learning information
- Makes the tiniest possible change to deceive the model



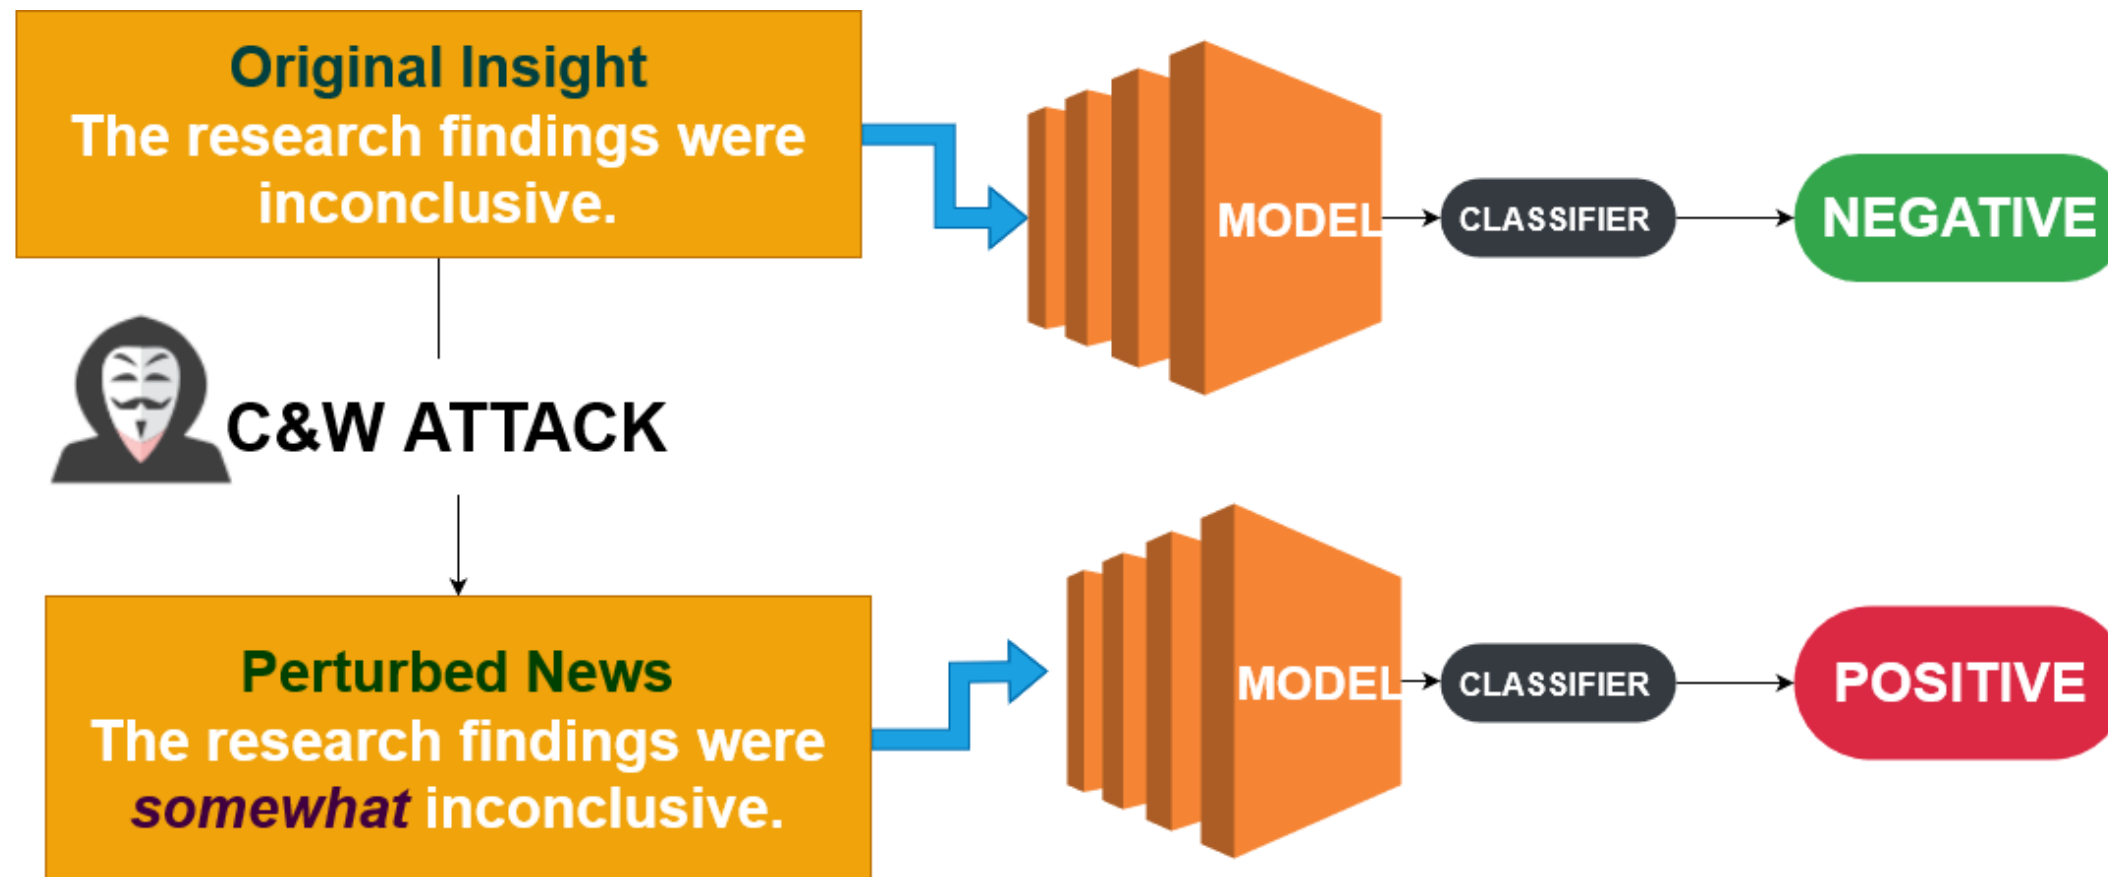
Projected Gradient Descent (PGD)

- More advanced than FGSM: it's iterative
- Tries to find the most effective disturbance



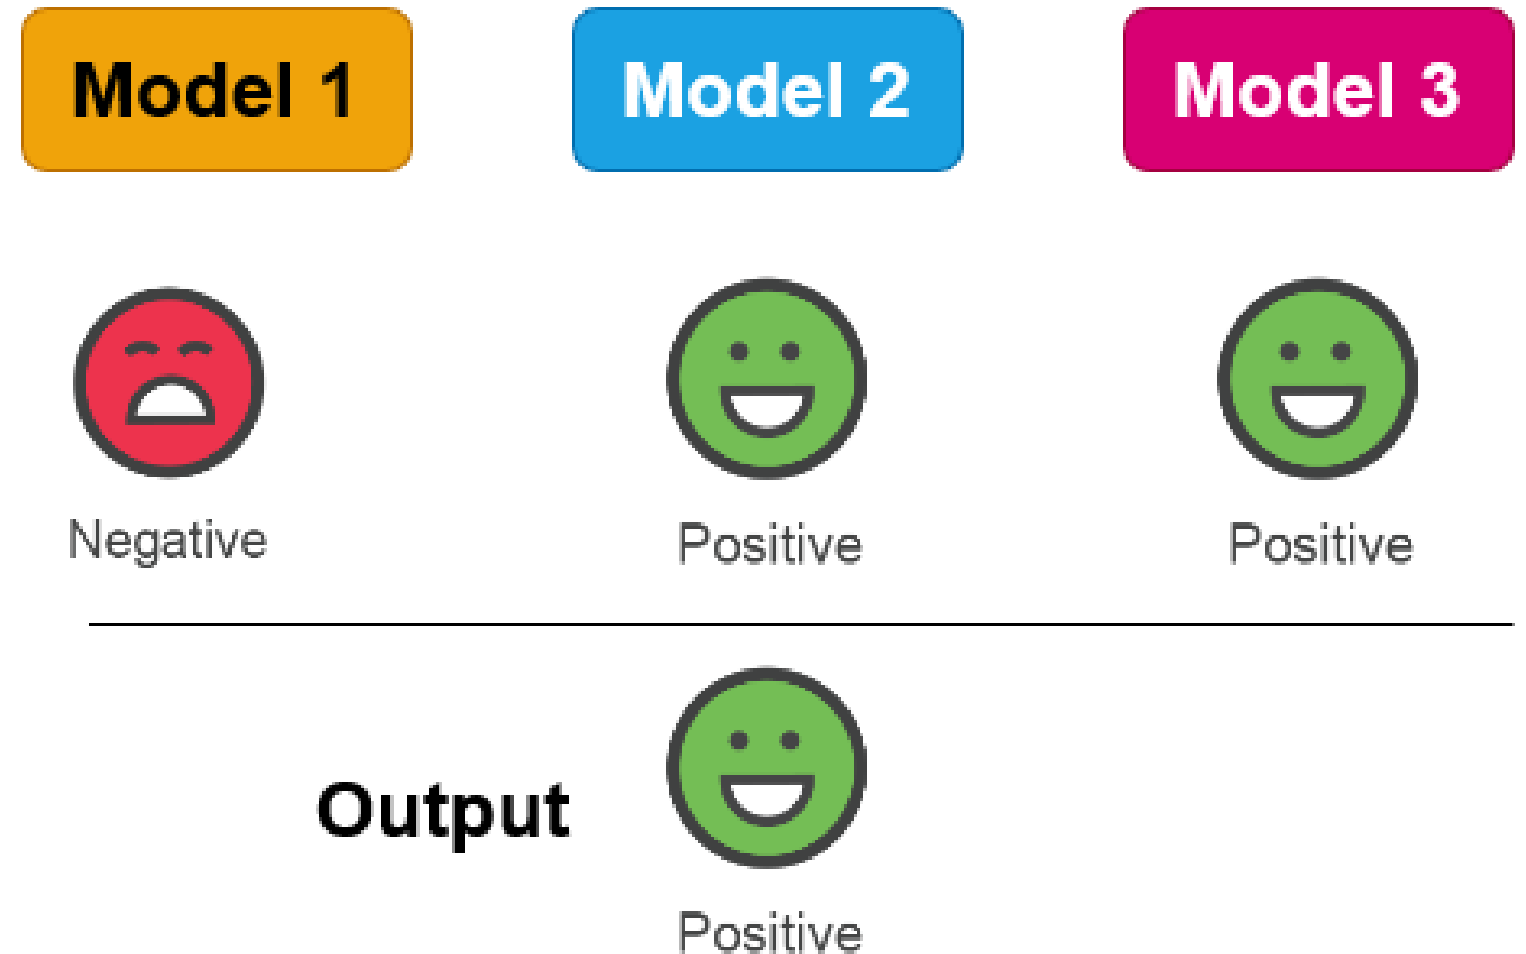
The Carlini & Wagner (C&W) attack

- Focuses on optimizing the loss function
- Not just about deceiving but about being undetectable




Building defenses: strategies

- **Model Ensembling:**
 - Use multiple models
- **Robust Data Augmentation:**
 - Data augmentation
- **Adversarial Training:**
 - Anticipate deception



Building defenses: tools & techniques

- **PyTorch's Robustness Toolbox:**
 - Strengthen text models
- **Gradient Masking:**
 - Add variety to training data to hide exploitable patterns
- **Regularization Techniques:**
 - Ensure model balance



¹ <https://adversarial-robustness-toolbox.readthedocs.io/en/latest/>,
<https://stock.adobe.com/ie/contributor/209161356/designer-s-circle>

Let's practice!

DEEP LEARNING FOR TEXT WITH PYTORCH

Wrap-up

DEEP LEARNING FOR TEXT WITH PYTORCH



Shubham Jain
Instructor

What you learned

- Chapter 1: Foundations of Text Processing
- Chapter 2: Text Classification Techniques
- Chapter 3: Text Generation Methods and Pre-trained Models
- Chapter 4: Advanced Deep Learning Topics

Key takeaways

- Encoding Text: one-hot, BoW, TF-IDF
- Deep Learning Models: CNN, RNN, GAN
- Advanced Techniques: Transformers & Attention
- Adversarial Attacks on Text Classification

Applied learning

- Implemented text classification models
- Built text generation models
- Used pre-trained models for text tasks
- Applied transfer learning

What's next?

- On DataCamp:
 - [Introduction to LLMs in Python](#)
 - [How to Train a LLM with PyTorch](#)
 - [Building a Transformer with PyTorch](#)
- Projects: text completion, chatbot text generation and sentiment analysis

Congratulations and Thank You!

DEEP LEARNING FOR TEXT WITH PYTORCH